

The Machine-Mediated Commerce Front

Where value actually accrues as agents — not people — become the buyers, and what the firm should advise, build and invest in across three horizons.

Prepared for the executive team of a multi-domain consulting firm

AT A GLANCE

17%

of shoppers are comfortable buying through AI

41%

of software organisations have MCP in production

88%

reported an AI-agent security incident this year

3–7 yrs

for the economic restructuring to play out



This deck answers five questions in sequence — each section closes one before the next opens

01	Is the shift real, and is it on time?	The shift is real. The popular timing and location of value are wrong.
02	Where does value actually accrue?	To scarce control points — not to protocols, which are commoditising.
03	When does it land, and in which sectors?	Enterprise and B2B first; consumer autonomy last. Six sectors ranked.
04	What do we do about it?	Three separate theses: for clients, for the firm, for external investment.
05	What would prove us wrong?	Five falsifiers and seven leading indicators, declared up front.

SO WHAT Each section is mutually exclusive and collectively exhaustive: no argument appears twice, and nothing material is left out.

The working thesis is confirmed — with one correction that changes where we invest first

SUBSTANTIALLY CONFIRMED

Machine-mediated commerce is real, accelerating and structurally significant

- Commerce is moving from human interfaces to machine-readable capabilities — high confidence
- Value accrues to controllers of intent, data, capability, identity, orchestration, settlement and assurance
- Open, foundation-governed standards are winning the protocol layer
- The restructuring is a three-to-seven year phenomenon, not a this-year event

ONE MATERIAL CORRECTION

The Horizon-1 front line is enterprise and the rails — not autonomous consumer shopping

- Value is in enterprise back-office and B2B workflows: procurement, service, finance, compliance
- And in the payment, identity and trust “rails” that make delegation safe
- Autonomous consumer checkout stalled in early 2026 on trust and conversion economics
- It is a Horizon-2/3 phenomenon — fund it as an option, not as a plan

SO WHAT Build around the durable control points — not around any single protocol, and not around the consumer-checkout narrative. MCP is now infrastructure, but it is one layer of a multi-protocol stack.

Consumers accept agents as shopping assistants far faster than as autonomous buyers

MARCH 2026

OpenAI scaled back in-chat Instant Checkout

- Only ~12 of Shopify's millions of merchants had gone live. The single most decisive datapoint of 2026.

WALMART

In-chat checkout converted ~3× worse

- The same shoppers sent to walmart.com converted at roughly three times the rate of completing in chat.

FEBRUARY 2026

eBay banned unauthorised buy-for-me bots

- Marketplaces are actively fencing off autonomous purchasing rather than courting it.


17%

of shoppers feel comfortable completing a purchase through AI

ChannelEngine Marketplace Shopping Behavior Report 2026. Trust — not model capability — is what is binding.

SO WHAT “Discover in AI, buy on site” is the durable pattern today. Advise clients to fund machine-readable capability and governance — not chatbot checkout.

Machine-mediated commerce nests in three scopes — and Horizon-1 value sits in the middle ring, not the inner one



Three propositions carry the whole argument — everything that follows is evidence for one of them

01

The shift is real, but the popular narrative is mis-timed

- Enterprise and B2B agentic workflows plus the payment-identity-trust rails are the Horizon-1 value
- Autonomous consumer shopping stalled on trust and conversion economics in early 2026
- Advise clients to invest in machine-readable capabilities and governance, not chatbot checkout

02

Value accrues to control points, not protocols

- MCP, A2A, ACP, UCP, AP2 and x402 are all converging on open, foundation governance
- That convergence will substantially commoditise them
- Defensible value sits in identity and delegated authority, differentiated data, orchestration, settlement and assurance

03

For the firm, this is a no-regret foundation play — now

- Build a proprietary maturity diagnostic, a trust/control library and protocol accelerators
- Stand up a managed agent-operations and assurance service
- Partner — do not build — on observability and payments; treat autonomous operators as a frontier bet

SO WHAT Sections 01–03 establish these propositions; sections 04–05 convert them into decisions and declare what would overturn them.

01

SECTION

The shift is real — but mis-timed


The infrastructure arrived in 2026. The trust did not.

IN THIS SECTION

- The interface shift: from attracting humans to serving agents
- Four milestones that made the infrastructure real
- Why announcements are not adoption
- Trust as the binding constraint on autonomy
- Where Horizon-1 value therefore sits
- The ten gates every autonomous use case must pass



The commercial question changes from “how do we attract a human to our interface?” to “how does a trusted agent transact with our capabilities?”



SO WHAT An 8,000-vendor martech stack is built on the assumption that buyers arrive at seller environments. Agents break that assumption — and with it, the discovery and SEO value pool.

In 2026 the infrastructure layer stopped being experimental — four milestones made it real

28 JULY 2026

MCP reached its stateless specification

- OAuth 2.1 / OIDC aligned; Resource Indicators (RFC 8707) required
- Enterprise-Managed Authorization now stable — Anthropic, Microsoft, Okta
- ~half a billion monthly Tier-1 SDK downloads; ~9,650 registry servers

9 APRIL 2026

A2A reached v1.0 — production-grade

- Signed Agent Cards and multi-tenancy for cross-boundary coordination
- 150+ organisations; Linux Foundation governed, absorbed IBM ACP
- Shipped in Vertex AI, Bedrock AgentCore, Azure AI Foundry

26 MAY 2026

Payment authority consolidated at FIDO

- Google's AP2 and Mastercard's Verifiable Intent both donated
- AP2 v0.2 adds "Human Not Present"; 60+ partners
- Payments TWG chaired by Mastercard and Visa

THROUGH 2026

Card networks shipped production agentic rails

- Visa Intelligent Commerce Connect unifies four protocols
- Mastercard Agent Pay live in Hong Kong (Mar) and Thailand (Apr)
- x402 became a Linux Foundation project, operational 14 July 2026

SO WHAT The plumbing is no longer the constraint. Any strategy that is still waiting for the standards to settle is now a year behind.

Announcements are not adoption — production reality is roughly half of what the headlines imply

41%

of surveyed software organisations have MCP in limited or broad production

Stacklok 2026 software report

8.5%

of MCP servers implement mandatory OAuth 2.1

The security floor is not being met

14.4%

of agents went live with full security and IT approval

Gravitee State of AI Agent Security 2026

88%

of organisations reported a confirmed or suspected AI-agent security incident

Gravitee 2026; healthcare 92.7%

THREE HEADLINE NUMBERS THAT DO NOT SURVIVE TRIANGULATION

- “MCP is 78% in production” is debunked — use ~41%
- x402 cumulative payment counts do not reconcile: Coinbase figures range ~75M to ~169M, on only \$24–50M cumulative value
- AgentCore “~200ms settlement” is the settlement step only — independent testing measured ~3–5s end to end

SO WHAT Read every vendor number as a pilot number until proven otherwise. Our credibility with clients depends on being the firm that corrects these, not the one that repeats them.

Trust, not technology, is what stops autonomy — and trust is an engineering programme, not a communications campaign

THE DEMAND-SIDE DEFICIT

Users will not delegate

- Only 17% of shoppers are comfortable buying through AI
- eBay banned unauthorised buy-for-me bots; Walmart saw ~3× worse conversion
- Liability for agent errors is unresolved across the FTC, the FCA and PSD3

THE SUPPLY-SIDE ATTACK SURFACE

Agents introduce new classes of attack

- Tool poisoning, descriptor poisoning, agent-in-the-middle, goal hijacking
- Documented in the OWASP MCP Top 10 and peer-reviewed research
- Real exploitation: CVE-2025-59528 in Flowise; a malicious MCP server impersonating Postmark

THE UNLOCK

Five controls turn autonomy on

- Just-in-time ephemeral scoped tokens
- Human-in-the-loop for every state-mutating operation
- Signed manifests and agent cards; allowlists
- Continuous monitoring and audit evidence

The strategic implication Every consulting engagement in Horizon 1 is, underneath, a trust engineering engagement. That is a services business — not a software business — and it is ours to take.

SO WHAT 88% incident rate plus 14.4% approval rate is the commercial opening: assurance is the highest willingness-to-pay category in the stack.

Because trust gates autonomy, Horizon-1 value sits where a human still approves — enterprise and B2B workflows

USE CASE	LIFECYCLE STAGE	EVIDENCE OF VALUE	AUTONOMY LEVEL	EXPOSURE
Procurement research and shortlisting	Discovery → evaluation	51% of B2B software buyers now begin in an AI chatbot; 69% chose a different vendor than planned	Recommendatory	Low
Quote-to-cash and PO matching	Authorisation → ordering	Up to 80% cycle-time reduction in PO processing (PwC); 40–60% cycle cuts (Salesforce)	Approval-based	Low
Credit risk memos	Evaluation → decision	20–60% productivity gain and 30% credit-turnaround improvement at a US bank (McKinsey)	Approval-based	Low
Financial crime and compliance	Service → renewal	FIS + Anthropic Financial Crimes AI Agent; BMO and Amalgamated Bank as early deployers	Governed, auditable	Low
Inventory monitoring and replenishment	Ordering → fulfilment	Agents monitor stock, compare terms and trigger purchases inside existing contracts	Bounded autonomy	Medium
Customer service and disruption recovery	Service	Malaysia Airlines “Mavis” end-to-end service agent; Navan travel rebooking	Bounded autonomy	Medium
Consumer discovery and comparison	Discovery → evaluation	High adoption — but the transaction step stalls on trust	Recommendatory	High at checkout

SO WHAT The four highlighted use cases combine quantified production outcomes with low exposure. They are where the first engagement should start — in every sector.

Ten gates must all pass before a use case may execute autonomously — and they cannot be averaged away

GATE 01 Identified actor and owner	GATE 02 Authenticated identity	GATE 03 Delegated authority	GATE 04 Data rights and permissions	GATE 05 Secure capability access
GATE 06 Transaction limits and controls	GATE 07 Human intervention and reversibility	GATE 08 Evaluation and monitoring	GATE 09 Audit evidence and incident response	GATE 10 Legal review and acceptable unit economics

Why this matters for the diagnostic A maturity score that averages these ten gates hides the single one that will stop the programme. We report the index only alongside its bottleneck and its failed gates.

SO WHAT This gate list is the spine of the proprietary diagnostic in section 04 — it is the part clients cannot self-assemble from open tooling.

02

SECTION

Value accrues to control points


Protocols are converging on open governance — and therefore commoditising.

IN THIS SECTION

- Six architectural layers, and why no protocol spans them
- The nine-protocol register, verified as of 18 August 2026
- Where MCP fits — and the four jobs it is wrong for
- The nine scarce control points where value accrues
- Who gains and who loses as value migrates
- Twelve business-model archetypes, scored



Agentic commerce needs six architectural layers — and no single protocol spans them



SO WHAT Treating MCP as “the commerce stack” is a category error. It owns layer 2 decisively and touches nothing else.

Nine mechanisms matter today — and every one of them is converging on open, foundation governance

PROTOCOL	LAYER	PURPOSE	STATUS	GOVERNANCE	ADOPTION EVIDENCE
MCP	2	Agent-to-tool and data connectivity	2026-07-28 stateless	LF Projects / AAIF	~½bn monthly SDK downloads; ~41% in production
A2A	3	Cross-vendor agent communication	v1.0, 9 Apr 2026	Linux Foundation	150+ orgs; Vertex, AgentCore, AI Foundry
ACP	5	Agent-to-merchant checkout	Beta, Apache 2.0	OpenAI / Stripe / Meta	PayPal, Agentforce — but scaled back Mar 2026
UCP	5	Full shopping journey	v2026-04-08	Google + Shopify	20+ endorsers; Walmart, Target, Etsy, Visa, Mastercard
AP2	5	Delegated payment mandates	v0.2, Apr 2026	FIDO Alliance	60+ partners; extends A2A and MCP
Verifiable Intent	5	Tamper-proof log of authorised actions	Contributed May 2026	FIDO (Mastercard)	Payments TWG chaired by Mastercard and Visa
Trusted Agent Protocol	5	Distinguish legitimate agents from bots	Open spec, Oct 2025	Visa-led	Cloudflare, Akamai, Adyen
x402	5	HTTP-native payments	Operational 14 Jul 2026	x402 Fdn / Linux Fdn	40 members — but volume heavily gamified
REST, EDI, RPA, iPaaS	2 / 4	Existing integration	Mature	Various	Often sufficient — and often the right answer

SO WHAT Standards position is worth having. Building a proprietary protocol is not — foundation governance will commoditise every one of these.

MCP is the right tool for capability exposure — and the wrong tool for four other jobs

Where MCP is the best fit

- Exposing enterprise tools and data to agents
- Internal agent-to-tool connectivity
- Anywhere an agent must discover and invoke a capability at runtime

Where MCP needs complements

- Cross-organisation coordination → A2A
- Commerce semantics → ACP or UCP
- Payment authority → AP2 / Verifiable Intent
- Settlement → x402 or card rails

Where existing APIs suffice


- Deterministic, high-volume B2B transactions
- EDI and REST flows with no goal interpretation
- Anything that does not meet the agentic bar

Where MCP adds avoidable complexity

- Simple request-response integrations
- Any integration with no agent in the loop
- Deterministic RPA dressed up as agentic

SO WHAT Gartner projects 75% of API-gateway vendors will ship MCP features by end-2026 — the product opportunity is protocol gateways, adapters and a governed semantic layer above MCP.

Connectivity will be commoditised — value accrues to nine scarce control points



SO WHAT Identity/authority and transaction orchestration are the two very-high-defensibility points — and both are exactly where the card networks and FIDO are consolidating.

Value migrates from human-facing discovery to machine-readable capability, identity and settlement

LIKELY TO GAIN

Payment networks

They own settlement, credentialing, fraud — and now the trust layer

Identity and trust providers

FIDO, Okta, Cloudflare, Akamai: agent authentication and delegated-authority credentials

Hyperscalers and cloud platforms

They own the agent runtime and the gateway

Well-integrated marketplaces

Structured catalogues become the moat — Shopify Agentic Storefronts, Google Shopping Graph

Data and knowledge providers

Differentiated inputs become directly monetisable per call

LIKELY TO LOSE

Pure discovery and SEO intermediaries

The 8,000-vendor martech stack assumes buyers arrive at seller environments

Online travel agencies

Disintermediation risk is highest in travel distribution — OTA channel cost is 15–25% of room revenue

Merchants with poor machine-readable data

Agents route around unparseable suppliers — the penalty is silent and total

Retail discovery advertising

Retail media must be reinvented for a reader that is not human

Anyone betting on a single protocol

Foundation governance is commoditising the layer they hoped to own

SO WHAT The losing side is not losing to a competitor — it is losing to a reader that never visits. That is why data and capability readiness is the no-regret first move.

Twelve archetypes exist — but only five combine high defensibility with a Horizon-1 start

ARCHETYPE	PAYER / BENEFICIARY	MONETISATION	DEFENSIBILITY	HORIZON
Identity, authority and trust provider	Enterprises and networks	Per-credential, subscription	Very high	H1–H3
Payment and settlement provider	Merchants and agents	Interchange, take-rate, FX	Very high	H1–H3
Evaluation, security and assurance provider	Enterprises	Subscription, services	High	H1–H2
Managed agent-operations provider	Enterprises	Managed fee, outcome-linked	High	H1–H2
Orchestration platform	Enterprises	Consumption	High	H1–H3
Enterprise agent provider	Enterprises	Licence, seat, usage	Medium	H1–H2
Merchant capability provider	Merchants	Platform fee	Medium	H1–H2
Data and knowledge provider	Agents and enterprises	Access licensing, x402 per call	High	H2–H3
Demand-side personal agent	Consumers	Subscription, take-rate	Low — assistant-dependent	H2–H3
Agent marketplace or broker	Both sides	Take-rate, listing	High if liquidity achieved	H3
Outcome-based autonomous operator	Enterprises	Success fee	High but capital-intensive	H3
Machine-resource marketplace	Machines and agents	Micro take-rate	Medium	H3

SO WHAT The five highlighted archetypes are where the firm should play. Four of the five are services businesses — which is the one place we start with an advantage.

03

SECTION

Where it lands, and when


Three horizons, six sectors, one ranking.

IN THIS SECTION

- The three-horizon interpretation — scenarios, not forecasts
- Six sectors ranked on Horizon-1 readiness
- Financial services: the most production-ready sector
- Retail: the most hyped and the least mature
- B2B and manufacturing: hardest, largest, data-gated
- Professional services: our own disruption



The restructuring is a three-to-seven year phenomenon — the next 18 months are foundations and bounded pilots



SO WHAT McKinsey’s \$3–5 trillion figure is a Horizon-3 scenario, not a Horizon-1 fact. Forecasts across sources range 35× — treat every one of them as a scenario to test.

Six sectors, one ranking: financial services is ready now; retail is the most hyped and the least mature

RANK	SECTOR	HORIZON-1 READINESS	WHERE THE VALUE IS NOW	THE CONTROL POINT TO HOLD
1	Financial services, payments, insurance	High — production	Credit, financial crime, claims, service	Governance-first deployment and the settlement rails
2	Manufacturing and B2B distribution	Medium-high	Procurement, quote-to-cash, PO matching	Machine-readable supplier and product data
3	Travel and hospitality	Medium-high	Booking, rebooking, disruption recovery	Real-time machine-readable availability; direct-booking recapture
4	Professional services	Medium	Research, drafting, review, diligence	Outcome-based pricing capability and delivery IP
5	Telecommunications and technology	Medium — operations only	Network operations, provisioning, NOC	Exposing connectivity, identity and billing as agent-callable
6	FMCG, consumer products, retail	Low on transaction	Discovery and comparison only	Catalogue normalisation, retail-media reinvention, identity linking

SO WHAT Readiness tracks one variable: whether a human already sits on the authorisation step. Sectors that already have approval chains are ready; sectors that rely on impulse are not.

Financial services is the most production-ready sector — and governance-first is the winning posture

70%

of banking leaders are deploying or exploring agentic AI

Sector survey, 2026

30–50%

workload reduction reported in banking operations

Cost-to-serve is the primary driver

20–60%

productivity gain on credit memos, with 30% faster turnaround

McKinsey, US bank

80%

faster insurance claims handling; near-instant settlements cited

Allianz

PLATFORM CONSOLIDATION

The core providers have already shipped

- Fiserv agentOS; FIS with Anthropic for financial crime
- Oracle Financial Services agentic platform, extended to corporate banking April 2026
- Wolters Kluwer: 44% of finance teams expected to use agentic AI in 2026

THE RAILS LAND-GRAB

Payments is the epicentre

- Visa Intelligent Commerce Connect; Mastercard Agent Pay and AP4M
- Mastercard's BVNK acquisition closed 3 August 2026 for up to \$1.8bn
- Nearly half of banks and insurers are creating dedicated agent-supervision roles

SO WHAT “Compliance as a design directive” is the sector’s own language for our assurance proposition — sell it in their words, not ours.

Retail is the most hyped and the least mature — the durable pattern is “discover in AI, buy on site”

THE REALITY Discovery adoption is high; transaction adoption is not ▪ A majority of consumers now use AI in shopping research ▪ Only 17% are comfortable completing the purchase ▪ OpenAI retreat; Walmart’s ~3× worse in-chat conversion	THE CONTROL POINTS Three things worth owning ▪ Machine-readable product data — catalogue normalisation is the moat ▪ Retail media reinvented for a non-human reader ▪ Loyalty and identity linking — UCP added Identity Linking in March 2026	THE LIKELY WINNERS Structured catalogues, not big brands ▪ Shopify Agentic Storefronts activate UCP and native MCP by default ▪ Google Shopping Graph plus Merchant Center plus Google Pay plus UCP ▪ This combination is the live closed-platform risk
THE SEQUENCING WE RECOMMEND Horizon 1 Data and catalogue readiness; agent discoverability. Horizon 2 Agent-enabled channels; participate in UCP and ACP without betting on either. Horizon 3 Autonomous replenishment.		

SO WHAT Never open a retail conversation with checkout. Open it with catalogue readiness — it is the prerequisite for every later move and it pays back regardless of how autonomy lands.

B2B is structurally the hardest and commercially the largest — and data quality is the competitive weapon

THE SIZE OF THE PRIZE

The largest single value pool in the research

- Gartner: by 2028, 90% of B2B buying will be AI-agent intermediated, pushing over \$15 trillion of B2B spend through agent exchanges
- Forrester: ~20% of B2B sellers will face agent-led quote negotiations during 2026
- G2: 51% of B2B software buyers now begin in an AI chatbot; 69% switched vendor on chatbot guidance

WHY IT IS HARD

Negotiated contracts, approval chains, EDI legacy

- Autonomous transaction is structurally hardest where terms are negotiated rather than listed
- Existing EDI and REST flows are often sufficient — and MCP would add avoidable complexity
- Oracle a Leader in the 2026 Gartner MQ for Source-to-Pay; SAP Catalog Optimization Agent

THE COMPETITIVE WEAPON

\$15M

average annual loss from poor data quality (Mirakl) — and it compounds under agentic procurement, because agents simply route around unparseable suppliers. The penalty is silent: you are not outbid, you are never surfaced.

THE AUTONOMY PROGRESSION

Earn each step — do not skip one

1. Research and market scanning
2. Shortlisting and recommendation
3. Approval-based ordering inside existing contracts
4. Bounded negotiation within pre-agreed parameters

SO WHAT Sell data readiness as revenue protection, not as an IT programme. In an agent-mediated market, unparseable is indistinguishable from absent.

Our own industry is the clearest business-model disruption — and outcome-based pricing is arriving more slowly than claimed

THE SIGNAL

McKinsey ~25% of global fees now outcome-linked; “Lilli” deployed at scale

BCG expects AI-tied work at ~40% of revenue in 2026, up from ~20% in 2024

Bain ~30% today, heading toward 50%

Accenture \$3.6bn AI bookings in FY2025 — up 120% year on year

EY and Deloitte have both published the accounting playbooks: EY on outcome-based pricing under ASC 606 (17 Feb 2026); Deloitte DART on agentic-AI pricing (4 Jun 2026)

Legal Harvey at 1,300+ organisations and 100,000+ lawyers; LexisNexis Protégé GA with 300+ workflows

THE COUNTER-SIGNAL

- The shift is sluggish. Even at McKinsey, roughly 75% of fees remain traditional billing.
- No public data quantifies what share of contracts are actually being renegotiated to outcomes — the headline percentages describe intent, not executed contracts.
- Outcome pricing transfers delivery risk to us. Without the assurance and evaluation capability from section 01, that is an uninsured transfer.
- The accounting treatment is only twelve months old — revenue recognition is a live constraint, not a solved one.

SO WHAT Treat outcome-based pricing as a capability to build deliberately over Horizon 2 — not as a forced pivot. Build the assurance capability first; it is what makes the risk transfer survivable.

04

SECTION

What we do about it

Three theses, kept deliberately separate.

IN THIS SECTION

- Why conflating the three theses is the common strategic error
- What we advise clients to fund — and to defer
- What the firm should build, and refuse to build
- Where the firm should invest externally
- The three-horizon roadmap, with two decision gates



Three investment theses answer three different questions — conflating them is the most common strategic error

A The client thesis

What should our clients fund?

- Owner: industry and client-account leadership
- Currency: client capital and client risk appetite
- Failure mode: recommending what we happen to sell

B The firm-capability thesis

What should we build ourselves?

- Owner: practice and capability leadership
- Currency: our own investment budget and partner time
- Failure mode: building what open tooling will commoditise

C The market thesis

Where should we invest, partner or acquire?

- Owner: corporate development and alliances
- Currency: balance sheet and alliance slots
- Failure mode: buying into a category we could simply partner in

SO WHAT Rank these three separately and review them separately. A category can be a strong client recommendation and a poor firm investment at the same time — observability is exactly that.

Advise clients to fund three no-regret foundations now, and to defer consumer autonomous checkout

RECOMMENDATION	CLASSIFICATION	HORIZON	RATIONALE
Machine-readable data and capability exposure	No-regret foundation	H1	Prerequisite for every downstream use case — agents route around unparseable data
Identity, delegated authority and consent architecture	No-regret foundation	H1	The binding gate for any autonomous execution; the regulatory direction of travel (FIDO, HM Treasury)
Agent evaluation, observability and security	No-regret foundation	H1	88% incident rate makes this a production blocker, not a maturity nicety
Bounded enterprise use cases with human-in-the-loop	Near-term priority	H1	Quantified ROI at low exposure — procurement, PO matching, financial crime, service
UCP / ACP participation and card-network agentic rails	Near-term priority	H2	Preserves channel access; hedge across both rather than picking a winner
Outcome-based and autonomous-operator models	Strategic option	H2–H3	High value but unproven economics and unresolved liability
Machine-native settlement (x402)	Monitor	H3	Volume is largely gamified today — watch for genuine adoption
Building proprietary protocols	Avoid	—	Protocols are commoditising under foundation governance

Default posture **BUILD** data, capability and identity foundations · **BUY or PARTNER** for observability, security and payment rails · **DEFER** consumer autonomous checkout

SO WHAT The three no-regret foundations pay back under every scenario in section 05 — including the ones where this thesis is wrong. That is what makes them the opening recommendation.

The firm should build proprietary diagnostics and assurance IP — and should refuse to build an observability platform

INVESTMENT	CLASSIFICATION	WHAT IT IS	COMMERCIAL MODEL
Proprietary agentic-commerce maturity diagnostic	No-regret foundation	Eight dimensions, ten gates, sector benchmarks; a 64–96 question bank	Recurring assessment revenue; lead-gen into transformation
Trust and control library plus assurance methodology	No-regret foundation	Six control domains mapped to a risk-control matrix	High willingness-to-pay given the 88% incident rate
Sector benchmarks and use-case libraries	Near-term priority	Six priority sectors, reusable across engagements	Differentiation and delivery speed
Protocol accelerators, reusable MCP assets and adapters	Near-term priority	Gateways, adapters and a governed semantic layer over MCP	Delivery speed plus recurring maintenance revenue
Evaluation and observability accelerators	Near-term priority	Assembled on OpenTelemetry, Confident AI and Arthur — not built	Partner-led and fast to stand up
Managed agent-operations and assurance service	Strategic option	Monitoring, evaluation and incident response as a service	Recurring-revenue annuity; sticky by design
Acquisition of niche agent-security or evaluation boutiques	Strategic option	A fragmented market with real talent scarcity	Talent and IP acceleration
A full proprietary observability platform	Avoid	Open tooling is moving too fast to out-run	Partner instead — the category will commoditise

SO WHAT The strategy is assemble-then-build: wrap best-of-breed open tooling inside proprietary diagnostic and control IP, and build only the executive-assessment, prioritisation and benchmark layers where the IP is genuinely defensible.

Externally, three categories are near-term priorities — and one is an explicit avoid

Identity, delegated authority and trust for agents	NEAR-TERM PRIORITY	Very high defensibility, standards-anchored at FIDO, high willingness-to-pay
Agent security, assurance and evaluation	NEAR-TERM PRIORITY	High willingness-to-pay in a fragmented market with credible acquisition targets
Data readiness and machine-readable catalogue services	NEAR-TERM PRIORITY	Clear and quantified pain; buyers already understand the problem
Protocol gateways, adapters and governed semantic layers	STRATEGIC OPTION	Durable precisely if standards fragment — a hedge against our own convergence call
Managed agent operations	STRATEGIC OPTION	Annuity revenue, but dependent on proving assurance demand first
Machine-native settlement and agent marketplaces	FRONTIER BET	Large addressable market but early, gamified volume and platform-dependent
Consumer personal-agent applications	AVOID / MONITOR	Low defensibility and wholly assistant-dependent

SO WHAT Scored on addressable value pool, willingness-to-pay, horizon, defensibility, switching costs, capital intensity, technical and regulatory risk, competitive intensity and platform dependence.

Two decision gates govern the roadmap — at 12 months on assurance demand, and at 36 months on outcome economics

Horizon 1 0–18 months

- Build the maturity diagnostic, trust/control library and sector use-case libraries
- Stand up an agent-security and assurance practice — partner on tooling
- Launch data/capability readiness and MCP-inventory assessments
- Establish alliances with a hyperscaler, a card network and an identity provider
- Recruit and retain agent-security and integration talent

KPIs

Diagnostic engagements sold · assurance ARR · alliance certifications

GATE AT 12 MONTHS — if assurance and diagnostic demand is proven, expand into managed operations

Horizon 2 18–36 months

- Scale managed agent-operations as recurring revenue
- Deliver cross-enterprise A2A workflow transformations
- Build outcome-based pricing capability internally, learning from the EY, Deloitte and BCG playbooks
- Develop sector reference architectures

KPIs

Managed-service ARR · outcome-based revenue share · repeat-client rate

GATE AT 36 MONTHS — if outcome-based economics prove out, shift more of the book to success-fee models

Horizon 3 3–7 years

- Position for agent-marketplace and machine-commerce advisory
- Develop dynamic-contracting IP
- Consider selective build or acquire in identity and assurance if the market consolidates

KPIs

Share of revenue from machine-commerce categories · IP licensing

SO WHAT Nothing in Horizon 2 is committed until the 12-month gate clears. The roadmap is deliberately optioned, because the falsifiers in section 05 are live.

05

SECTION

What would prove us wrong

The falsifiers and indicators, declared before the fact.

IN THIS SECTION

- Five developments that would falsify or weaken the thesis
- Seven leading indicators, and the two that matter most
- The assumptions this research is built on
- Six decisions for the executive team



Five developments would falsify or materially weaken this thesis — we are declaring them before the fact

01

Consumer autonomous checkout recovers

If conversion reaches parity with redirect, the consumer sub-thesis we downgrade here would need upgrading — and Horizon-1 spend should shift toward consumer channels.

02

Protocols fragment permanently

If ACP and UCP never converge, integration costs rise and the fragmentation counter-thesis wins. Gateways and adapters become more valuable, not less.

03

Rails commoditise to near-zero margin

If incumbents capture all the value, there is no new value pool and no consulting annuity — only a one-off integration wave.

04

Regulatory clampdown stalls delegated authority

EU high-risk classification of agentic commerce, or PSD3 SCA rules that block autonomous payment, would freeze Horizon 2.

05

Machine-native settlement never becomes genuine

Roughly half of x402 activity was wash or gamified at the March 2026 snapshot, with only ~\$28,000 a day genuine. If that never converts, Horizon 3 shrinks.

SO WHAT None of these five is currently triggered. Two of them — 01 and 05 — are the ones we should instrument now, because they would change the investment case rather than merely delay it.

Seven indicators tell us whether we are right — and two of them matter more than the other five

MCP production-adoption rate	~41% today, with OAuth 2.1 at ~8.5%	
ACP / UCP convergence signals	The one genuinely contested layer in the stack	
x402 genuine, wash-filtered volume	The falsifier for the whole Horizon-3 case	PRIORITY
Card-network agentic transaction counts in live APAC markets	Hong Kong and Thailand are the leading edge	
Consumer autonomous-checkout conversion versus redirect	The single most consequential number in this deck	PRIORITY
FIDO agentic-standards ratification	Determines how fast delegated authority becomes routine	
PSD3 / PSR application date and any agent-specific provisions	Earliest ~H2 2027, potentially slipping to 2028	

SO WHAT Instrument the two priority indicators as a standing quarterly review item. They are the two datapoints that would most change the thesis — everything else is confirmation.

Four assumptions shape every number in this deck — they should be challenged before the conclusions are

ASSUMPTION 01

The agentic bar is defined and enforced

“Agentic” requires goal interpretation, decision-making, tool use, adaptation or delegated action. Deterministic RPA and EDI are excluded unless they meet that bar.

ASSUMPTION 02

Announcements are not adoption

Pilots are not production. Every adoption figure in this deck distinguishes the two, and vendor claims are triangulated with independent evidence or flagged where they cannot be.

ASSUMPTION 03

Market sizes are scenarios, not predictions

Forecasts range 35× across sources — from eMarketer’s ~\$144bn to McKinsey’s ~\$5tn by 2030 — driven purely by definitional boundaries.

ASSUMPTION 04

Two variables carry most of the sensitivity

The findings are most sensitive to whether consumer autonomous checkout conversion recovers, and whether machine-native settlement volume becomes genuine. Both are flagged as falsifiers.

SO WHAT If the room disagrees with a conclusion, take the disagreement back to one of these four assumptions. That is where the real argument lives.

Six decisions for the executive team

BELIEVE

The front is real and consequential — but the next 18 months of value are in enterprise and B2B workflows and the identity-payment-trust rails, not autonomous consumer shopping.

BUILD

A proprietary maturity diagnostic, a trust/control and assurance library, sector use-case libraries and benchmarks, and protocol accelerators — assembling, not rebuilding, open tooling.

INVEST IN

Agent identity and trust, security and assurance, and data readiness near-term; managed operations and protocol gateways as options; machine-native settlement as a frontier bet.

ADVISE

Fund machine-readable capability, identity and delegated authority, and agent assurance as no-regret foundations. Run bounded pilots with a human in the loop. Defer consumer autonomous checkout.

PARTNER FOR

Observability and evaluation, payment rails and identity — via alliances with one hyperscaler, one card network and one identity provider.

MONITOR

The seven leading indicators — above all whether consumer autonomous-checkout conversion recovers, and whether machine-native settlement volume becomes genuine.